

The Joint Production of Price Efficiency and Heavy-Tailed Volatility in Minimal Agent-Based Markets: Added Rationality Converts the Sign Leak but Cannot Close It

Abstract

Agent-based models of financial markets reproduce the stylized facts of asset returns piecemeal: successes are exhibited and failed attempts rarely reported. We score minimal flow-based markets on a fixed ten-event battery calibrated so that real equity data passes all ten and geometric Brownian motion exactly three. A market of fundamentalists, chartists, volatility-targeting investors, and market makers attains five of ten; the missing events concern long memory, time-irreversibility, and the information content of price signs. Three pre-registered experiments test whether added rationality closes the gap: a trading anticipator of the forecastable flow, the fixed point of mutual anticipation iterated to depth five, and a price-setting maker that provably absorbs that flow into the price level. All three are refuted, and a registered attribution program establishes what the interventions do instead: each converts a genuine nonlinear multi-day sign leak into one-lag microstructure reversal — absorbing most of it (whitened information 0.0197 to 0.0036 bits at exact absorption), never closing it (significant in every configuration) — and pays with the amplitude facts: heavy tails (kurtosis $8.8 \rightarrow 3.0$), the leverage effect ($-0.125 \rightarrow -0.008$), and volatility clustering are deleted together. In this model class, sign efficiency and the heavy-tailed, asymmetric volatility regularities are jointly produced by the same state-dependent flows — scoped by measurement to flow-generated amplitude structure and robust to liquidity re-equilibration and threshold perturbation. Cross-venue evidence — the leverage effect ordered across equities (-0.04), foreign exchange (≈ 0), and cryptocurrencies ($+0.03$) — is consistent with the mechanism. All results are pre-registered, validated on synthetic ground truth, and regenerable from the repository.

Keywords: stylized facts; agent-based models; market efficiency; price impact; leverage effect; pre-registered research.

JEL classification: G14; G17; C63; C52.

1 Introduction

Daily asset returns exhibit two properties simultaneously. Their signs are nearly unforecastable from public information, and their amplitudes display the heavy-tailed, asymmetric regularities documented across markets and decades: excess kurtosis, volatility clustering, the leverage effect, and gain/loss tail asymmetry (Cont, 2001). Throughout the paper we refer to the first property as *sign efficiency* and to the second group — the amplitude-side stylized facts — as the *amplitude facts*. This paper asks whether a minimal market mechanism can generate both properties at once, and answers by refutation: in flow-based minimal markets, sign efficiency and the amplitude facts are jointly produced by the same flows, and three pre-registered attempts

to obtain them separately — rationality added as a trading agent, as the fixed point of mutual anticipation, and as a price-setting rule — all fail against refutation criteria declared in advance. The refutations are the result: they compose into a structural statement about what this model class cannot do, measured on a fixed instrument.

Three research questions organize the analysis. First, which of the stylized facts can a minimal flow-based market reproduce when all of them are scored jointly on a fixed, pre-calibrated instrument rather than exhibited selectively? Second, can bounded rationality — introduced as a trading agent, iterated toward its fixed point, or embedded in the pricing rule itself — supply the facts such a market misses, in particular the near-absence of sign information in prices? Third, is the mechanism that answers the first two questions consistent with how the amplitude facts vary across real trading venues with different flow structures?

The questions sit in a literature with a shared target list: the stylized facts that any candidate mechanism for price formation should reproduce (Cont, 2001). Agent-based modeling took up the challenge early: heterogeneous-agent markets with fundamentalists and chartists generate fat tails and clustered volatility from agent interaction rather than from exogenous news (Lux and Marchesi, 1999), and a substantial literature has since mapped which ingredients produce which facts (LeBaron, 2006). Two gaps in that literature motivate the design here. First, scoring is usually piecemeal: a model is exhibited that matches two or three facts, with the remaining facts unexamined, so the literature accumulates existence proofs rather than a ledger of what a given mechanism class can and cannot reproduce. Second, negative results are structurally underproduced: when an added ingredient fails to improve a model, the failure is rarely reported, almost never pre-registered, and therefore carries no evidential weight. We address both gaps with a fixed measurement instrument and a pre-registered refutation chain.

The paper makes five contributions. *First*, it introduces a ten-event battery of binary stylized-fact tests, computed identically on real and simulated series, calibrated at both ends of its range: real S&P 500 index data (SPY) passes 10/10 events while geometric Brownian motion passes exactly 3/10, so intermediate scores are interpretable (Section 4.1). *Second*, it establishes, through three successive pre-registered experiments with pre-specified refutation criteria and parameter budgets, that the best flow-only configuration’s score of 5/10 is invariant to added rationality of three distinct kinds — a trading anticipator, the iterated fixed point of mutual anticipation, and a price-setting market maker — and that the refutations compose into a structural claim: in this model class the residual sign information and the amplitude facts are generated by the same state-dependent flows (Section 6). *Third*, it develops a methodology that makes such negative results evidential: every estimator and every tested mechanism must pass a size and power validation against synthetic data-generating processes with known properties before any evaluation data is read, so a refutation is a measurement rather than an absence of result (Section 4). *Fourth*, it subjects its own central experimental reading to a registered attribution and robustness program — six experiments answering three independent referee reports, with budgets and decision rules fixed before any run: a matched-strength design shows the anticipation-depth axis was an effective-strength axis; an AR(1)-whitened decomposition of the sign-information statistic, licensed by a size and power gate on an analytic AR(1) world, shows the raw rises are induced reversal while the genuine nonlinear leak is monotonically absorbed but never closed; and the composed claim is stress-tested against an exogenous amplitude source, liquidity re-equilibration, and threshold perturbation (Section 6.5). The program overturns one natural reading of the raw evidence and sharpens the surviving mechanism statement to conversion without closure. *Fifth*, it reports a cross-venue

measurement of the leverage effect with a single estimator battery — equities -0.04 , foreign exchange statistically zero, cryptocurrencies $+0.03$ in point estimate — whose ordering is consistent with the flow-structure mechanism, together with an explicit power analysis of what the cross-venue data can and cannot yet certify (Section 7).

The headline mechanism result was reached in two stages, and both are reported. The minimal market’s central failure is a *sign leak* — residual forecastability of next-day return signs, defined precisely in Section 6, that real index data does not show. The three experiments refute the hypothesis that added rationality removes it: one layer of anticipation leaves the battery score unchanged; iterating anticipation five layers deep contracts the model-space forecastable flow by a factor of four while the raw E9 statistic moves the opposite way (median direction bits $0.0184 \rightarrow 0.0200 \rightarrow 0.0257$); and a maker that removes the forecastable component’s price impact exactly — verified to the identity on a synthetic world — deletes the fat tails (kurtosis $8.8 \rightarrow 3.0$), the leverage effect ($-0.125 \rightarrow -0.008$), and volatility clustering, with the raw statistic again higher. A registered attribution program (Section 6.5) then establishes what the raw statistic was measuring. The rises are intervention-induced one-lag reversal, scored as sign information by a statistic whose feature set conditions on the return sign: under AR(1) whitening every rise reverses (positive on 0 of 8 seeds in all four arms), while the genuine nonlinear sign leak falls monotonically with intervention strength — from 0.0197 bits in the control to 0.0036 under exact absorption — and remains statistically significant in every configuration (on 39 of 40 configuration–seed pairs). Added rationality converts the sign leak into microstructure reversal; it absorbs most of the leak, closes none of it, and pays for the conversion with the amplitude facts. Sign efficiency and the amplitude facts are not separate dials in this model class.

The remainder of the paper is organized as follows. Section 2 reviews the related literature. Section 3 describes the data. Section 4 presents the measurement methodology: the event battery, the size/power validation protocol, and the pre-registration design. Section 5 specifies the base market model and the three rationality mechanisms. Section 6 reports the results of the ablation study, the three experiments, and the registered attribution and robustness program. Section 7 presents the cross-venue evidence. Section 8 states the scope of the claims and the limitations, and Section 9 concludes. Appendices A–C contain exact threshold definitions, the multi-venue calibration audit, and robustness material.

2 Related work

2.1 Stylized facts of asset returns

Heavy tails in speculative price changes were documented by Mandelbrot (1963); Cont (2001) fixed the canonical list of stylized facts — heavy tails, volatility clustering, the leverage effect, near-absence of linear return autocorrelation, gain/loss asymmetry — that this paper’s battery operationalizes, and our contribution to that program is calibration at both ends (real data 10/10, GBM exactly 3/10) so that partial scores are interpretable. In reduced form, volatility clustering is captured by the ARCH family (Engle, 1982; Bollerslev, 1986), with asymmetric extensions such as GJR-GARCH (Glosten et al., 1993) adding the leverage channel. The battery deliberately includes an event this family fails: amplitude correlations in real data decay hyperbolically rather than exponentially (Ding et al., 1993), a message sharpened by the rough-volatility literature (Gatheral et al., 2018). The measurement program behind this paper

replicated a Hurst exponent $H \approx 0.10$ on 16 years of SPY range-based volatility, and event E4 is calibrated so that exponential-memory processes (GJR-GARCH) fail it — which is exactly the event no flow configuration and no rationality layer ever produced, consistent with rough volatility’s message that the clock’s memory is structural, not flow-generated. Time-reversal asymmetry, the target of event E8, is a documented stylized fact in its own right (Zumbach, 2009). Outside equities, the stylized-facts program has been extended to cryptocurrency markets (Urquhart, 2016; Zhang et al., 2018), whose distinct flow structure Section 7 exploits.

2.2 Heterogeneous-agent models

The agent-based literature descending from Lux and Marchesi (1999) and surveyed by LeBaron (2000, 2006) and Hommes (2006) established that fundamentalist–chartist interaction generates fat tails and clustering, whether through switching between predictor types, herding in the communication structure (Cont and Bouchaud, 2000), or the destabilizing feedback of technical strategies on the price they condition on (Farmer and Joshi, 2002; Chiarella et al., 2009). Our ablation reproduces the classical result (and its fragility: chartists without market makers destroy efficiency, scoring 1/10) but differs in asking the complement — which facts a flow-based market *cannot* reproduce at any rationality depth, established by pre-registered refutation rather than by exhibition.

Two adjacent traditions ask closely related questions with different instruments. In the heterogeneous-beliefs models of Brock and Hommes (1997, 1998), the sophistication of agents’ expectations is a bifurcation parameter of a deterministic pricing map: making agents more forward-looking changes the dynamics qualitatively rather than monotonically improving them. Our K -ladder is a stochastic, capital-bounded analogue of that exercise, with a cautionary coda: its apparent depth effect dissolves under a matched-strength attribution — away from the capital caps the ladder telescopes to a single anticipator whose effective strength, not its depth, drives the measured statistic (Section 6.5). The minority-game literature (Challet and Zhang, 1997; Challet et al., 2004) and the intermittent agent-based markets of Giardina and Bouchaud (2003) derive stylized facts from the statistical mechanics of adaptive agents competing over a scarce resource, with realistic regimes appearing near phase boundaries; our battery differs in scoring a fixed list of calibrated events rather than exhibiting a regime, but the joint-production finding is consistent with that literature’s lesson that efficiency and excess volatility are properties of one collective state, not separate features.

2.3 Market efficiency and the sign channel

The battery’s E9 event operationalizes the oldest target in the field. Samuelson (1965) proved that properly anticipated prices fluctuate randomly, and Fama (1970) organized the empirical program around informational efficiency. A classical argument holds that when information is costly, full informational efficiency cannot be an equilibrium: some residual predictability must survive to compensate the agents who remove it. The joint-production result is a mechanical relative of that argument with the cost relocated: in this model class the residual sign information is not a fee paid to arbitrageurs but a by-product of the same state-dependent flows that generate the volatility structure, and the three experiments show it cannot be traded, iterated, or quoted away without taking the amplitude facts with it.

2.4 The leverage effect

The classical explanations of the equity leverage effect are financial leverage in the firm’s capital structure (Christie, 1982) and volatility feedback through the equity risk premium (Bekaert and Wu, 2000); Bouchaud et al. (2001) measured the effect’s retarded temporal structure across individual stocks and indices and modeled it directly at the price-process level. Both classical mechanisms are equity-specific, which is what makes the cross-venue comparison of Section 7 informative: a currency is not a levered claim on productive assets and carries no equity risk premium, so whatever return–volatility asymmetry survives outside equities must have another source.

2.5 Order flow, price impact, and de-leveraging feedback

The price-impact literature of Bouchaud and coauthors (Bouchaud et al., 2009, 2018) underlies the paper’s central object: the idea that order flow is the proximate mover of prices and that liquidity provision digests predictable flow. Our quote-skewing maker is a minimal formalization of perfect digestion — and Experiment III’s finding that exact absorption deletes the volatility dynamics along with the leak is, in that vocabulary, a statement that the market’s amplitude structure *is* its undigested flow. The de-leveraging spiral at the model’s core is a minimal form of the funding-liquidity feedback of Brunnermeier and Pedersen (2009) and of the procyclical intermediary leverage documented by Adrian and Shin (2010): a volatility rise forces balance-sheet contraction, and the contraction moves prices. Volatility-targeting mandates, which hard-wire this feedback into a mechanical trading rule, are the model’s concrete instance of the mechanism.

2.6 Reflexivity and endogeneity

Filimonov and Sornette (2012) measured near-critical reflexivity in markets via the Hawkes branching ratio, and Hardiman et al. (2013) showed the estimate is sensitive to kernel specification, finding near-critical values throughout two decades of data. The measurement program behind this paper found the raw result replicates (branching ratio 0.68 across 48 US assets) but collapses to 0.25 — a no-self-excitation null — once events are deformed by the volatility clock, so most apparent reflexivity at the daily horizon is volatility clustering; the same deformation-first discipline shapes the battery’s E8 event.

2.7 Multiple testing and pre-registration in finance

Harvey et al. (2016) documented the multiple-testing problem behind the factor literature and argued that published findings should clear substantially higher evidential hurdles; Bailey and López de Prado (2014) and Bailey et al. (2017) supplied the deflated Sharpe ratio and the probability of backtest overfitting as concrete corrections. The protocol of Section 4.3 — pre-registered hypotheses, bounded search budgets, deterministic tie-breaks, and refutation criteria honored when they fire — is that program’s counterpart for mechanism research rather than performance research: it is what makes the negative results of Section 6 reportable as measurements.

3 Data and descriptive statistics

All series are daily. Four data sets enter the paper: the US equity calibration series, the multi-index audit set, the foreign-exchange panel, and the cryptocurrency panel. All were obtained from Yahoo Finance as daily open/high/low/close (OHLC) bars and are cached in the repository’s data layer; every screen described below runs inside the loaders, so a failing screen aborts the study rather than silently degrading it.

3.1 Equities

The battery’s calibration anchor is SPY, the S&P 500 index ETF, over 2010–2026 (close-only, approximately 4,100 trading days). The event thresholds themselves descend from a prior measurement program on ~ 48 liquid US equities and ETFs over the same span (Section 4.1); panels in that program require 95% coverage per instrument and forward-fill gaps of at most three days. The multi-index audit set (Appendix B) comprises six further index ETFs on a fixed, registered list: QQQ, DIA, and IWM (US, 2010–2026) and three locally-listed non-US index trackers — 1306.T (TOPIX), EXW1.DE (EURO STOXX 50), and 2800.HK (Hang Seng). The cross-venue equity cohort of Section 7 consists of four universes of locally-listed large-capitalization equities (US, Japan, Europe, Asia-EM), each in a single time-zone block, from the research program’s transfer study.

3.2 Foreign exchange

The FX panel contains 13 liquid crosses — twelve majors plus USDMXN as a pre-declared emerging-market stress case — over 2010–2026, the same span as the equity universes. Yahoo FX bars occasionally carry degenerate ranges (high equal to low, or the range pinned to $|\text{close} - \text{open}|$), which would invalidate range-based variance estimation (Garman and Klass, 1980); a screen measured the fraction of days with a real intraday range per pair before the study design was locked. All 13 pairs pass the pre-declared 95% bar (measured real-range fractions 99.88–100.00%), and none were dropped.

3.3 Cryptocurrencies

The registered cryptocurrency panel contains the 10 largest coins with full-period daily history (BTC, ETH, XRP, LTC, BCH, ADA, DOGE, LINK, XLM, ETC) over 2017-11-09 to 2026-06-04, the span for which major-coin daily data exist; every coin passes the same real-range screen at 100% of days. A registered robustness widening to 17 majors (adding TRX, XMR, EOS, NEO, DASH, ZEC, BAT; ATOM excluded by a pre-declared history rule) is used in the power analysis of Section 7 and leaves the registered 10-coin panel untouched.

3.4 Descriptive statistics

Table 1 reports universe-level descriptive statistics for the six panels used in the cross-venue analysis. Two regularities in the table organize the paper. First, the kurtosis of volatility-deformed returns sits near the Gaussian value in every venue while raw kurtosis does not: conditionally on their own volatility path, daily returns are close to Gaussian everywhere, which is what the battery’s event E6 tests in simulation. Second, the leverage effect changes sign across venues — negative in every equity universe, indistinguishable from zero in FX, positive

Panel	N	Span	kurtosis	kurtosis (deformed)	leverage effect
US equities	49	2010–2026	13.1	3.4	−0.047
Japan equities	29	2010–2026	8.9	3.3	−0.039
Europe equities	36	2010–2026	11.3	3.6	−0.046
Asia-EM equities	29	2010–2026	8.4	3.8	−0.030
FX crosses	13	2010–2026	7.6	2.7	+0.005
Cryptocurrencies	10	2017–2026	16.6	4.5	+0.031

Table 1. Universe-level descriptive statistics. N is the number of instruments on each registered list (per-instrument coverage screens may drop instruments from individual analyses). Statistics are the measurement program’s pooled universe-level estimates: kurtosis of daily returns, kurtosis of volatility-deformed returns $r_t/\hat{\sigma}_t$, and the leverage effect $\text{corr}(r_t, \text{RV}_{t+1..10})$. The deformed-kurtosis column collapses toward the Gaussian in every venue; the leverage column changes sign across venues.

in point estimate in cryptocurrencies — and that ordering is the external-validity evidence analyzed in Section 7. The full statistic vector for the battery’s calibration anchor (SPY) appears in Table 10 of Appendix A.

4 Methodology

4.1 The ten-event battery

The battery consists of ten binary events, computed identically on real data and on simulated output, with pass criteria fixed before any agent-based run. Table 2 lists them; Appendix A gives the exact statistic definitions and thresholds as implemented, together with the full statistic vector at both calibration ends. The events are close-only versions of regularities measured and gate-validated in a prior program on ~ 48 liquid US equities and ETFs (2010–2026), so the thresholds are empirical ranges, not conventions.

Calibration pins both ends of the scale: real SPY (close-only) scores 10/10; a pure GBM scores exactly 3/10, passing only E1, E6 and E9 — Gaussian worlds are efficient, trivially one-clock, and carry no sign information. Everything between 3 and 10 is what a candidate mechanism must supply. Two scoping remarks apply. The thresholds are empirical ranges from a US program, and a registered audit on six further indices (Table 11 in Appendix B) measures how far the joint anchor extends: DIA also scores 10/10, QQQ and IWM 8/10, while three non-US index ETFs score only 5–7/10 — the battery is a US-calibrated instrument, and its function in this paper is comparative — every configuration is scored on the same events with the same seeds — not to certify absolute realism for any venue. Simulated scores carry a seed-bootstrap standard error — the eight evaluation seeds are resampled with replacement 2000 times and the majority-vote score recomputed per resample; the SD of the resampled scores is reported as $\pm$ in the experiment tables — and across every configuration in this paper that SE is at most 0.5 events. What protects the comparisons beyond the SEs is that they are paired (all configurations share seeds 100–107) and that every conclusion below rests on a pre-registered refutation criterion, not on a one-event margin.

#	Event	Pass criterion
E1	price efficiency	$AC_1(r) \in [-0.15, 0.05]$ (allows the real index’s mild daily reversal)
E2	fat tails	$kurtosis(r) \in [4.5, 40]$
E3	vol clustering	$AC_1(r) \geq 0.12$ and $\text{mean } AC(r , \text{lags } 5\text{--}20) \geq 0.05$
E4	long-memory clock	8-week volatility-clock level autocorrelation above threshold
E5	leverage effect	$\text{mean corr}(r_t, RV_{t+1..10}) \leq -0.03$
E6	one-clock Gaussianization	$kurtosis(r/\sigma_{5d}) \leq 5$
E7	clock jumps	AR(1) clock-innovation skew beyond the $\log\text{-}\chi^2$ noise floor
E8	arrow of time in the coupling	entropy production significant in r and reduced $\geq 25\%$ by vol deformation
E9	no sign information	direction bits $\leq$ shuffle null (95%)
E10	tail asymmetry	gain/loss tail asymmetry at 2.5σ

Table 2. The ten-event battery (final calibrated form; exact estimator definitions and thresholds in Appendix A). E1, E4, E7 and E10 were recalibrated against real data before any ablation was run, under a gate requiring real SPY to score 10/10 and geometric Brownian motion exactly 3/10; notably, a GJR-GARCH process (Glosten et al., 1993) fails E4 — exponential memory cannot mimic long memory. E9 measures the mutual information between three public tape features and the next-day return sign with a discrete plug-in estimator against a permutation null.

4.2 Size and power validation on synthetic ground truth

Every estimator in the program passes a validation gate before touching evaluation data: a test on synthetic data where the answer is known, establishing both size (on a world where the effect is absent, the test does not fire) and power (on a world where the effect is present by construction, it does) (Figure 1). The battery’s estimators were licensed this way before the first agent-based run, and — the step specific to this paper — so was every *tested mechanism*: the anticipator was shown to profit and damp flow impact on a synthetic world (Section 5.2), the fixed-point stack to contract forecastable flow (Section 5.3), the quote-skewing maker to absorb impact while leaving the flow untouched (Section 5.4). This is what makes the refutations evidential: when a licensed mechanism fails to move the battery, the failure is a property of the market, not of a broken implementation. At the time of writing the continuous-integration suite runs all 40 verification gates (the numbered gates X1–X36 plus four unnumbered platform gates) on every commit.

4.3 Pre-registration, budgets, and tie-breaks

“Pre-registered” here means: the design document — hypotheses, budgets, refutation criteria — is committed to the repository before any evaluation run, and the version-control history is the timestamp; there is no external registry. Each experiment fixed, before any evaluation run: hypotheses (with one designated primary), the full parameter set, the evaluation protocol (8 seeds, IDs 100–107, $T = 6000$, majority vote per event, identical across all three experiments so ablations are paired seed-for-seed), a refutation (“kill”) criterion, and a bounded tuning budget — one grid pass on disjoint tuning seeds (900–903), permitted only under a pre-declared




Figure 1. The validation pattern. Every estimator and every tested mechanism must, before any evaluation data is read, both detect a planted effect on a synthetic world where the effect exists (power) and remain silent on a world where it does not (size). The examples are this paper’s: the fixed-point operator’s contraction gate and the quote-skew maker’s absorption gate.

trigger, with a deterministic pre-declared tie-break (“weakest intervention wins”). Amendments were permitted only before evaluation data was read, and each is recorded in the registration documents with its trigger. Two features of this protocol proved consequential. First, the refutation criteria fired three times and were honored three times. Second, the tie-break rule turned the tuning passes themselves into evidence: in all three experiments the battery’s preferred setting was the weakest intervention available — the market scores best when the rational agent barely acts — which is the structural claim visible in the selection pattern before it is visible in the scores.

4.4 Reproducibility discipline

Three practices matter for the trustworthiness of small numeric differences. (1) *Byte-identity pins*: every new mechanism ships behind a flag, and a gate asserts that flag-off output is byte-identical to the pre-change simulator, so the calibration (SPY 10/10, GBM 3/10) and all previously published rows are protected across the entire experimental sequence; Experiment II’s $K=1$ reproduces Experiment I’s published configuration exactly, and Experiment III’s $\lambda_Q=0$ reproduces both. (2) *Causality gates*: tampering with future inputs must leave the trade or quote prefix bit-identical (and the tail must differ, proving the mechanism actually reads the tape). (3) *Architecture-portable assertions*: an early version of the pins asserted absolute byte hashes of floating-point output, which held on one CPU architecture and failed on another by units in the last place; the discipline that survives is to assert in-process equivalence and architecture-stable invariants rather than absolute hashes.

Separately from this paper’s claims: performance claims elsewhere in the same research program are governed by a trial ledger and reported as deflated Sharpe ratios (Bailey and López de Prado, 2014) beside a combinatorially cross-validated probability of backtest overfitting (Bailey et al., 2017); the program’s standing rule is that its PBO of 0.45 is restated beside any Sharpe claim. This paper makes mechanism claims, not performance claims, but the same

accounting culture (Harvey et al., 2016) — count the trials, publish the failures — produced it.

5 The model

5.1 The base market

The base market is deliberately small: aggregated flows, about a dozen parameters, linear impact. Log price p_t moves by

$$r_t = \lambda [D_t^{\text{fund}} + D_t^{\text{chart}} + D_t^{\text{vol}} + D_t^{\text{mm}} + D_t^{\text{noise}}], \quad (1)$$

with price-impact coefficient λ ; $D_t^{\text{fund}} = k_F (V_t - p_t)$ for a random-walk log fundamental V ; $D_t^{\text{chart}} = k_C \tanh(m_t/s)$ for an EWMA trend signal m with saturation scale s ; $D_t^{\text{vol}} = k_V (L_t - L_{t-1})$ for volatility-targeters (hereafter vol-targeters) running the leverage rule $L = \min(L_{\max}, \sigma^*/\hat{\sigma})$ against an EWMA variance estimate $\hat{\sigma}^2$; a one-lag market-maker flow $D_t^{\text{mm}} = -k_M r_{t-1}$; and i.i.d. noise. The volatility-targeting flow is the reflexive ingredient: a vol spike forces de-leveraging (selling), which moves price, which raises vol — the minimal form of the leverage-constraint feedbacks of Brunnermeier and Pedersen (2009) and Adrian and Shin (2010). The market-maker ingredient was added at the calibration stage, before any ablation was scored, when tuning showed that vol-targeting flows structurally leak momentum — their de-leveraging is serially correlated — and that in reality such flow is absorbed by liquidity provision.

Parameters were set by hand once, on the full configuration, to produce realistic volatility levels and pass E1–E3 only, and were not tuned per configuration or per event; ablations zero out flows and change nothing else. Eight seeds per configuration, $T = 6000$; an event passes if a majority of seeds pass.

5.2 The anticipatory agent (Experiment I)

The vol-targeters’ flow is mechanical and forecastable from public information alone: their mandate ($L = \min(L_{\max}, \sigma^*/\hat{\sigma})$) and their estimator (an EWMA of squared returns at a known speed) are both computable from the price tape. The anticipatory agent A is given exactly that information and nothing else. At each step it reconstructs $\hat{\sigma}_t^2$ from the return history, forecasts the integrated future mechanical flow under the one belief that defines vol-targeting — volatility reverts to the targeters’ own target, so leverage ends at $L_{\text{eq}} = \min(L_{\max}, 1)$ —

$$\hat{F}_t = k_V \cdot \text{mean}_{\text{cohorts}}(L_{\text{eq}} - L_t), \quad (2)$$

holds inventory proportional to the forecast, capped by capital, $I_t^* = \text{clip}(k_A \hat{F}_t, \pm \text{cap}_A)$, and trades the inventory change plus execution noise of scale s_A . The unwind is automatic: as vol reverts and the targeters re-lever, the forecast decays and the anticipator sells into exactly that mechanical bid. Three parameters (k_A, cap_A, s_A) — response, capital, noise — frozen before any battery run.

Three gates license the mechanism before any score is read (the pattern is described in Section 4.2): (a) with the agent disabled the simulator reproduces the original market byte-for-byte, protecting the SPY-10/GBM-3 calibration; (b) the forecast is causal — tampering with future returns leaves the trade prefix unchanged; (c) on a deterministic synthetic world with a perfectly predictable mechanical flow, the anticipator profits *and* damps the flow’s price

impact, from 1.00 to 0.75 per unit of arriving flow. The gates establish that the agent functions as designed. The pre-registered refutation criterion: if the frozen configuration does not exceed 5/10 on the evaluation seeds, the expectation hypothesis is refuted.

The parameter protocol matters for what follows. The first-shot parameters $(k_A, \text{cap}_A, s_A) = (0.5, 0.02, 0.002)$ scored 4/10 on tuning seeds disjoint from evaluation, so the single pre-registered grid pass was spent: its best score was 5/10, tied across all six $k_A = 0.25$ combinations plus one weak $k_A = 0.5$ setting — flat wherever the anticipator is weak. The pre-declared deterministic tie-break (lexicographically smallest, i.e. the weakest anticipator) selected $(0.25, 0.01, 0.001)$. That the selection lands on “trade as little as possible” was recorded, at registration time, as already being evidence about the hypothesis.

5.3 The fixed point of mutual anticipation (Experiment II)

The fixed-point stack makes each layer’s unwind part of what is anticipated. Layer k front-runs the *residual* forecastable total flow — the mechanical flow plus the deterministic future unwind of the $k - 1$ layers beneath it (time subscripts suppressed):

$$\text{resid}_0 = \hat{F}, \quad J_k = \text{clip}(k_A \text{resid}_{k-1}, \pm \text{cap}_A), \quad \text{resid}_k = \text{resid}_{k-1} - J_k, \quad I_K = \sum_{k=1}^K J_k. \quad (3)$$

Away from the caps this telescopes: $I_K = (1 - (1 - k_A)^K) \hat{F}$, and the model-implied forecastable residual of total flow contracts geometrically, $\text{resid}_K / \hat{F} = (1 - k_A)^K \rightarrow 0$. $K = 0$ is the flow-only control; $K = 1$ is byte-for-byte Experiment I’s single layer; $K \rightarrow \infty$ is the rational-expectations limit of this operator class. The budget was fixed before any run: $K = 5$ exactly, no convergence-based stopping, evaluation on the same eight seeds, and a contingent six-candidate tuning pass permitted only on a regression below the control’s 5/10, with the same weakest-anticipation tie-break.

One amendment was made at the gate stage, before any battery run, and it illustrates what the gates are for: the first formulation capped the *stack*’s total inventory at a single agent’s cap_A , and the pre-specified mechanism gate failed it — under a shared cap the $K=5$ stack holds exactly the $K=1$ inventory through every large episode and the measured forecastable-flow contraction disappears. A total-cap stack cannot express the hypothesis it exists to test. The licensed formulation caps each layer separately.

With per-layer caps the mechanism gate passes: on the deterministic synthetic world the projection of realized total flow onto the mechanical flow contracts as $\beta_K = 1.000 \rightarrow 0.750 \rightarrow 0.237$ at $K = 0 \rightarrow 1 \rightarrow 5$, matching the $(1 - k_A)^K$ theory line exactly. The operator does what the rational-expectations argument says it should. The question is whether the market obeys the argument.

5.4 The price-setting market maker (Experiment III)

Experiments I and II eliminate every agent that *trades* against the forecastable flow: an anticipator’s inventory is driven by the same slow public state, so its trades re-leak what they absorb. One hypothesis survives that diagnosis: *price-setting* rationality — a market maker who does not trade but quotes, pre-adjusting the price by the expected impact of the forecastable flow so that the information is absorbed into the price *level* rather than left in the return. The maker reuses the same causal public-state forecast $\hat{F}_t$ of Eq. (2), holds no inventory, posts no

	E1	E2	E3	E4	E5	E6	E7	E8	E9	E10	
SPY (real)	■	■	■	■	■	■	■	■	■	■	10/10
G	■	□	□	□	□	■	□	□	■	□	3/10
F	■	□	□	□	□	■	□	□	■	□	3/10
FC	□	□	□	□	□	■	□	□	□	□	1/10
FV	□	■	□	□	■	■	■	□	□	■	5/10
FCV	□	■	□	□	■	■	■	□	□	■	5/10
FCVM	■	■	□	□	■	■	□	□	□	■	5/10
FCVMH	■	□	□	□	■	■	□	□	□	■	4/10

Figure 2. Battery scorecard: real SPY against the ablation ladder (filled = pass). G = noise only (GBM); F adds fundamentalists; C chartists; V volatility-targeters; M market makers; H multi-horizon heterogeneity. No flow configuration exceeds 5/10, and the five events missing from the best configuration (FCVM) are E3, E4, E7, E8, E9.

flow, and draws no noise:

$$q_t = \lambda_Q \lambda \hat{F}_t, \quad r_t = \lambda D_t + (q_t - q_{t-1}), \quad (4)$$

with D_t the total flow in brackets in Eq. (1) and skew strength $\lambda_Q \in \{1.0, 0.5\}$ fixed in advance ($\lambda_Q = 0$ is byte-identical to the unmodified simulator). Because the maker’s forecast state *is* the targeters’ public state, the quote revision telescopes exactly against the mechanical flow $f_t^{\text{mech}} = D_t^{\text{vol}}$, $q_t - q_{t-1} = -\lambda_Q \lambda f_t^{\text{mech}}$: at full skew the mechanical flow’s entire price impact is absorbed into the level by algebraic identity, not by tuning, and the realized return keeps only the unforecastable flow surprise. Nothing is left that could re-leak — no inventory, no unwind, no cap, no execution noise. The flow itself still executes untouched; only where its impact lands (level versus return) changes.

The mechanism gate verifies this to the identity on the deterministic synthetic world: at $\lambda_Q = 1$ the leak estimator $\text{corr}(f_t^{\text{mech}}, r_{t+1})$ collapses from +0.66 to −0.06 while the mechanical-flow series itself is essentially unchanged — absorption into the level, not suppression of the flow. The registration declared this the third and final attempt on the expectation line and included a closure clause: if the primary hypothesis is refuted at this budget, the 5/10 ceiling is declared structural for this model class and the line closes.

6 Results

6.1 The ablation ladder and the five-event ceiling

Figure 2 shows the ladder. Noise alone (G) and fundamentalists alone (F) score the Gaussian 3/10. Chartists without liquidity provision are destabilizing: FC scores 1/10, because the momentum leak destroys even efficiency. The vol-targeters are the decisive ingredient: FV

	$AC_1(r)$	kurtosis	$AC_1(r)$	leverage	E9 bits
SPY (real)	-0.103	15.1	0.31	-0.113	0.0029
GBM (G)	0.001	3.0	-0.001	0.007	0.0001
FCVM (best flow-only)	0.036	8.8	0.38	-0.125	0.0184

Table 3. Key medians. FCVM reproduces efficiency, tails, clustering amplitude and the leverage effect at realistic magnitudes; what it cannot do is keep the sign channel silent (E9: 0.0184 significant direction bits against SPY’s 0.0029) or generate long memory and the coupling arrow.

Config	Score	Failed events	E9 bits (median)
FCVM (control)	5/10 $\pm$ 0.18	E3 E4 E7 E8 E9	0.0184
FCVM+A (hypothesis)	5/10 $\pm$ 0.03	E3 E4 E7 E8 E9	0.0200
FV+A	5/10 $\pm$ 0.10	E1 E3 E4 E8 E9	0.0158
F+A (no targeters)	3/10 $\pm$ 0.16	= F exactly	0.0002

Table 4. Experiment I ablation (8 evaluation seeds, majority vote; $\pm$ is the seed-bootstrap SE of the score, Section 4.1). Adding the anticipator changes none of the targeted quantities: the failure set is identical and the direction bits do not fall.

jumps to 5/10, generating the asymmetric amplitude facts — fat tails (E2), the leverage effect (E5), clock jumps (E7), crash asymmetry (E10) — because the de-leveraging spiral is the only flow that reacts to volatility and reacts in one direction only. Adding market makers (FCVM) restores efficiency (E1) at the cost of E7, holding the score at 5/10: efficiency and the amplitude facts are contributed by *different* agents. Multi-horizon heterogeneity (FCVMH) fails both ways, dropping to 4/10.

The ceiling is 5/10 and the five missing events constitute a diagnosis (Table 3). FCVM fails E3 (slow volatility clustering), E4 (long-memory clock), E7 (clock jumps, marginally), E8 (the arrow in the coupling), and E9 (information-free signs). The E9 failure is central to the remainder of the paper, so we give it a name used throughout: the *sign leak* — a significant excess of E9’s statistic, the *direction bits* $I(\text{tape features; sign } r_{t+1})$, over its shuffle null. FCVM leaks 0.0184 bits where real SPY shows 0.0029. The leak has a known source: the vol-targeters’ de-/re-leveraging drift is serially correlated and forecastable from the tape, and a one-lag market maker cannot absorb a multi-day drift. The natural hypothesis — the one the three experiments test — is that the missing ingredient is *expectation*: rational agents who anticipate the predictable flow should arbitrage the sign leak away and perhaps stretch the volatility episodes into long memory.

6.2 Experiment I: one anticipator

Table 4 shows the outcome. The refutation criterion is met: FCVM+A scores 5/10 with a failure set identical to the control’s. The events expectation was hypothesized to supply did not move — the direction bits did not move toward zero (medians $0.0184 \rightarrow 0.0200$; the paired comparison, extended to 32 seeds under a registered fixed budget, is statistically flat — the anticipator above control on 19 of 32 seeds, two-sided Wilcoxon $p = 0.57$ — so the licensed reading at one layer is “unchanged,” not “reduced”), and the long-memory events stayed flat. The interaction hypothesis — that the anticipator can act only through the flow it fronts — held:

$k_A = 0.25$	5	5	5	5	5	5
$k_A = 0.5$	5	4	4	4*	3	3
$k_A = 1$	4	4	4	4	4	4
	0.01/0.001	0.01/0.002	0.02/0.001	0.02/0.002	0.05/0.001	0.05/0.002
	cap _A / s _A					

Figure 3. The Experiment I tuning grid (tuning seeds 900–903, disjoint from evaluation): battery score by anticipator response k_A , capital cap_A , and execution-noise scale s_A . Darker cells score higher. The score is non-increasing in every strengthening direction; the tied best settings are the weakest anticipators. The heavy box marks the tie-break selection (the frozen evaluation setting); the asterisk marks the pre-registered first-shot setting.

without vol-targeters the anticipator contributes nothing ($F+A = F = 3/10$) — expectation requires a predictable flow to trade against; having one, however, repairs none of the missing events.

The reason the attempt fails is itself the finding. The tuning grid is monotone (Figure 3): the harder the anticipator trades, the worse the market scores. Efficiency breaks (its unwind adds reversal: AC_1 falls from -0.09 at the frozen setting toward -0.35 at $k_A = 1$), the tails erode (it damps the very spiral that generates them, kurtosis $8.8 \rightarrow 4.6$ from the control to the grid’s most capitalized settings), and the raw E9 statistic rises (up to 0.039 bits at full strength $k_A = 1$) — a rise the attribution program of Section 6.5 identifies as the induced reversal itself, scored as sign information. The full grid — per-setting scores and statistics — is exported with the archived results.

The mechanism is structural, not parametric: the anticipator’s inventory is a function of the same slow public volatility state as the flow it front-runs, so its own future unwind is exactly as forecastable as what it absorbs. One layer of expectation does not silence the sign channel; the whitened decomposition of Section 6.5 shows it absorbing part of the genuine nonlinear leak (whitened bits $0.0197 \rightarrow 0.0112$) while introducing the one-lag reversal that keeps the raw statistic level — and eroding the amplitude facts wherever it trades harder. The refutation names its own successor: if one anticipator re-leaks because its inventory is public-state-driven, perhaps information-free prices require anticipators who anticipate *each other* — the fixed point of mutual anticipation.

6.3 Experiment II: the fixed point of mutual anticipation

Table 5 and Figure 4 give the result, which is the paper’s central experiment. The refutation criterion is met again, and this time the two diagnostics move in *opposite directions*: as K goes $0 \rightarrow 1 \rightarrow 5$, the model-space forecastable flow falls $1.000 \rightarrow 0.750 \rightarrow 0.237$ while the raw E9 statistic rises $0.0184 \rightarrow 0.0200 \rightarrow 0.0257$ bits in the medians, with $K=5$ above the control on 7 of 8 paired seeds and on 28 of 32 in the registered extension of Section 6.5. (The one-layer step inside that rise is a point estimate within seed noise — the registered

Config	Score	Failed events	E9 bits (median)
$K=0$ (FCVM, control)	$5/10 \pm 0.18$	E3 E4 E7 E8 E9	0.0184
$K=1$ (Exp. I layer, byte-identical)	$5/10 \pm 0.03$	E3 E4 E7 E8 E9	0.0200
$K=5$ (frozen carry-over)	$3/10 \pm 0.49$	E1 E3 E4 E5 E7 E8 E9	0.0257
$K=5$ (tuned, contingent pass)	$5/10 \pm 0.00$	E3 E4 E7 E8 E9	0.0168

Table 5. Experiment II: the K -ladder. The raw E9 statistic rises in K — the depth-five rise holds on 7 of 8 paired seeds and stands in a registered 32-seed extension, while the one-layer step ($K=0 \rightarrow 1$) is statistically flat at 32 seeds (Wilcoxon $p = 0.57$) — and the model-space forecastable flow contracts toward zero; the frozen $K=5$ stack is a two-event regression (it breaks efficiency and removes the leverage effect). Section 6.5 attributes the rise: at matched effective strength, depth five leaks *less* than depth one, and the raw rise is induced one-lag reversal, not sign information.

32-seed extension is statistically flat, medians 0.0195 versus 0.0194, Wilcoxon $p = 0.57$.) The frozen $K=5$ stack is moreover a two-event regression: its aggregate unwind adds reversal ($AC_1: -0.075 \rightarrow -0.247$, breaking E1) and it absorbs the very spiral that generates the leverage effect (E5 fails, $-0.095 \rightarrow -0.030$), while tails erode (kurtosis $6.8 \rightarrow 5.3$) and crash asymmetry collapses (tail ratio $44.8 \rightarrow 4.3$). Rational front-running at depth does not merely fail to supply the missing events; at carry-over strength it removes amplitude facts the market already had.

The regression triggered the pre-registered contingent pass: six candidates on disjoint tuning seeds returned a best score of 5/10, tied across the three weakest settings, and the pre-declared tie-break again selected the weakest — $k_A = 0.05$, an effective stack strength of $1 - 0.95^5 = 0.226$, *weaker than Experiment I’s single frozen layer* (0.25). Evaluated once on the evaluation seeds, the selected setting scores 5/10, failing exactly the control’s five events. The selection pattern of Experiment I repeats one level up: the tuning criterion rewards weaker anticipation at every opportunity, never stronger.

The algebra explains the failure sharply, and the attribution program of Section 6.5 completes it. Away from the caps, the K -layer stack is equivalent to *one* anticipator of strength $1 - (1 - k_A)^K$ (the telescoped sum in Eq. (3)): mutual anticipation of this operator class deepens into a stronger single anticipator, and Experiment I already measured that stronger is worse. The equivalence is not only algebraic but measured: holding effective strength at 0.25 exactly, a depth-five stack leaks *less* than the depth-one layer (medians 0.0176 versus 0.0194, Wilcoxon $p = 0.0028$), so the K -axis of Figure 4 is an effective-strength axis — the frozen stack ran at strength 0.763 against the single layer’s 0.25 — and what strength buys is not sign information but induced reversal (Section 6.5). The contraction that defines the fixed point is an open-loop property, licensed by the gate against an exogenous tape; it does not survive closed-loop equilibrium, because every layer’s inventory is a function of the same slow public volatility state, so the stack’s own price impact regenerates forecastable structure in proportion to its strength. Model-space forecastable flow down, market-space conversion up: stacking boundedly-capitalized flow traders cannot reach the rational-expectations limit — it only strengthens the single trader the stack telescopes to.

6.4 Experiment III: absorption removes the amplitude facts, not the leak

Table 6 and Figure 5 give the sharpest refutation of the three. The full-skew maker removes the vol-targeters’ forecastable impact exactly — and the sign channel still does not close. The raw statistic stays flat at half skew (0.0180 versus 0.0184) and rises at full skew (0.0237, higher




Figure 4. Experiment II in both spaces. Solid curve and filled circles (left axis): the forecastable-flow fraction β_K measured on the deterministic synthetic world, falling exactly as $(1 - k_A)^K$ — $1.000 \rightarrow 0.750 \rightarrow 0.237$. Dashed line and open squares (right axis): the raw E9 statistic measured in the simulated market at the same K (open circles: per-seed values — 32 seeds at $K \in \{0, 1\}$ from the registered extension, the 8 evaluation seeds at $K=5$), medians rising $0.0184 \rightarrow 0.0200 \rightarrow 0.0257$; the $K=5$ rise holds on 7 of 8 paired seeds, while the one-layer step is flat at 32 seeds. The operator removes forecastable flow in model space, and open-loop contraction does not survive closed-loop equilibrium — but the rise on the right axis is not re-created sign information: the whitened decomposition of Section 6.5 attributes it to intervention-induced one-lag reversal, and at matched effective strength depth five leaks less than depth one.

than the control on 7 of 8 seeds, and on 26 of 32 in the registered extension). The whitened decomposition of Section 6.5 states what each number means: exact absorption removes most of the genuine nonlinear leak — whitened bits fall $0.0197 \rightarrow 0.0036$ — without ever silencing it, while the absorption leaves the one-lag maker’s reversion dominating the return (efficiency breaks, $AC_1: +0.04 \rightarrow -0.26$) and the E9 statistic, whose feature set conditions on the return sign, scores that induced reversal as sign information. The leak was never the vol-targeters’ drift alone — Experiment I’s original attribution is overturned: a significant nonlinear residual survives exact removal of the spiral’s impact, so it lives in the remaining state-dependent flows.

The second pre-registered hypothesis — that the maker would neutralize only the forecastable *mean* of the flow, leaving the volatility dynamics intact — is also refuted at the theory case, and its refutation is the paper’s central fact: in this market the forecastable flow *is* the generator of the amplitude facts. Full absorption removes the fat tails (kurtosis $8.8 \rightarrow 3.0$), the leverage effect ($-0.125 \rightarrow -0.008$), crash asymmetry (tail ratio $40 \rightarrow 1.1$) and volatility clustering ($AC_1(|r|)$ $0.38 \rightarrow 0.06$) simultaneously. Half skew regresses as well ($3/10$), so the contingent pass fired: the six-candidate grid returned a uniform $5/10$ for every $\lambda_Q \leq 0.30$, and the pre-declared weakest-skew tie-break picked $\lambda_Q = 0.05$ — the same selection pattern a third time, now for a pricing rule. Evaluated once on the evaluation seeds, the selected setting scores $5/10$, failing exactly the control’s five events. Across all three experiments and all three agent classes, the

Config	Score	Failed events	E9 bits (median)
FCVM (control)	$5/10 \pm 0.18$	E3 E4 E7 E8 E9	0.0184
FCVM+Q($\lambda_Q=1.0$) (theory case)	$1/10 \pm 0.34$	all but E6	0.0237
FCVM+Q($\lambda_Q=0.5$)	$3/10 \pm 0.03$	E1 E2 E3 E4 E7 E8 E9	0.0180
FCVM+Q($\lambda_Q=0.05$) (tuned, contingent)	$5/10 \pm 0.00$	E3 E4 E7 E8 E9	0.0176

Table 6. Experiment III. Full skew removes the vol-targeters’ forecastable impact exactly, and the sign channel does not close — the raw E9 statistic rises (0.0237, higher than the control on 7 of 8 seeds; induced reversal per Section 6.5) — while the amplitude facts disappear.

tuning criterion selects the weakest rational intervention available.

Figure 6 collects the E9 evidence across all three experiments.

6.5 Attribution and robustness: the registered referee program

The results above report the battery’s raw statistics. Three independent referee reports on a draft of this paper accepted the refutations but challenged their reading, identifying six specific threats: a depth/strength confound in the K -ladder (away from the caps the stack telescopes to one anticipator of strength $1 - (1 - k_A)^K$, so Figure 4 varies strength, not only depth); an autocorrelation confound in E9 itself (its feature set conditions on $\text{sign}(r_t)$, and the correlation between $\text{AC}_1(r)$ and the raw statistic is -0.89 across the archived tuning grid, so induced one-lag reversal is mechanically scored as sign information); eight-seed power behind the two raw rises; the identification of amplitude structure with forecastable flow by construction; a market never re-equilibrated around the added rationality; and integer scores on marginal thresholds. All six were answered under a single registration — six experiments, R1–R6, with budgets (272 battery runs and 64 auxiliary simulations, held exactly) and decision rules fixed before any run — and the estimator the program needed was licensed first: gate X36 verifies that the whitened-MI estimator reproduces the battery’s E9 block exactly, matches the closed-form sign information of an analytic AR(1) world, removes exactly the linear channel there, and preserves a planted nonlinear sign channel. The registration’s reporting standard is applied below: each experiment ends in a verdict about which claim survives.

Matched strength (R1). Holding effective strength at 0.25 exactly — $K=1$ with $k_A = 0.25$ against $K=5$ with per-layer $k_A = 1 - (1 - 0.25)^{1/5} \approx 0.0559$ — on 32 seeds, depth five leaks *less* than depth one: medians 0.0176 versus 0.0194, two-sided Wilcoxon $p = 0.0028$, depth five higher on 6 of 32 seeds. The caps bind only for the single strong layer, and the realized effective strengths (0.2500 at depth five, 0.2237 at depth one) make the comparison conservative: the depth-five arm ran slightly stronger and still leaked less. Across the archived tuning grid, the Spearman correlation of anticipator strength with the raw statistic is 0.892. The verdict: the published K -axis was a strength axis, “deeper anticipation grows the sign leak” is dead as a depth attribution, and what strength grows is identified by R2.

Attribution of the E9 statistic (R2). Table 7 gives the decomposition and Figure 7 shows it. The control’s leak is genuinely nonlinear: AR(1) whitening leaves it essentially unchanged (0.0197 whitened against 0.0184 raw, significant on 8 of 8 seeds, $\hat{\varphi} = +0.036$). Every rationality layer *converts* it: whitened bits fall monotonically with intervention strength,




Figure 5. Amplitude-fact medians versus quote skew λ_Q (8 evaluation seeds). Full absorption of the forecastable flow’s price impact removes the fat tails (kurtosis $8.8 \rightarrow 3.0$), the leverage effect ($-0.125 \rightarrow -0.008$) and volatility clustering ($\text{AC}_1(|r|)$ $0.38 \rightarrow 0.06$) — while the raw E9 statistic (panel d; open circles are seeds) does not close: flat at half skew, higher at full skew, the rise being induced reversal (Section 6.5).

$0.0197 \rightarrow 0.0112$ ($K=1$) $\rightarrow 0.0038$ ($K=5$) and 0.0036 ($\lambda_Q=1$), while the sign-alone component rises $0.0101 \rightarrow 0.0220$ and 0.0206 and $\text{AC}_1(r)$ is driven to -0.25 . Under whitening the raw rises of Tables 5 and 6 do not merely vanish — they reverse (the whitened difference against the control is positive on 0 of 8 seeds, $p = 0.0078$, in all four arms), the sign-alone component carries more than the whole raw rise (share 2.21 at $K=5$, 1.84 at $\lambda_Q=1$), and the conditional information given $\text{sign}(r_t)$ falls in every arm ($p = 0.0078$). Per the registered decision rule, the reading that the market re-creates sign information in price space is *withdrawn*. Two claims are licensed in its place: the leak never closes — whitened bits remain significant in every configuration, on 8 of 8 seeds everywhere except the half-skew maker’s 7 of 8, reaching no lower than 0.0033 bits (half skew) and 0.0036 at exact absorption — and the interventions absorb most of the genuine nonlinear leak while manufacturing the one-lag reversal that the raw statistic scores.

Power (R3). The two raw rises stand at 32 seeds: the frozen $K=5$ stack at 0.0258 against the control’s 0.0195, higher on 28 of 32 seeds, Wilcoxon $p < 10^{-4}$; full skew at 0.0239, higher on 26 of 32, $p < 10^{-4}$. The power concern is answered in the direction of the published point estimates: the raw statistic robustly rises. Composed with R2, the rises are retained as well-powered measurements of induced efficiency breakage — reversal scored by a sign-conditioned statistic — not of sign-information creation.




Figure 6. The raw E9 statistic across every rationality configuration (open circles: evaluation seeds; bars: medians). No configuration approaches the information-free regime; real SPY (dotted, 0.0029) sits far below all of them. Trading against the flow, iterating that trade to depth five, and absorbing the flow’s impact into the price level all leave the sign channel open; what varies across configurations is the split between genuine nonlinear leak and induced reversal (Section 6.5), not the distance to silence.

A second amplitude source (R4). In the model as published, the only amplitude generator is the vol-targeting spiral, which is also the only forecastable flow, so “absorption deletes the amplitude facts” risks being true by construction. A registered probe adds an exogenous source: Student- $t(3)$ fundamental innovations, entering through the fundamentalist flow and hence outside the span of the maker’s forecast. The scale was fixed by one registered calibration shot (the variance-matching scale won, median kurtosis 8.82 against the control’s 8.79; the ladder is non-monotone — larger t -scales *lower* total kurtosis by drowning the spiral). The control FCVM-T3 scores 5/10, failing exactly FCVM’s five events; under full absorption, FCVM-T3+Q1.0 scores 1/10 with kurtosis 3.06 and the raw statistic at 0.0229 bits. The $t(3)$ tails do not survive absorption because the fundamentalist channel ($k_F = 0.15$) spreads a fundamental jump over roughly seven trading days — it is a structural tail filter. The verdict scopes the composed claim: joint production is structural to *flow-generated* amplitude structure. Whether sign efficiency and the amplitude facts are separately obtainable given a genuinely absorption-proof exogenous amplitude source is untested by this design and would require a different injection point (heavy-tailed noise-trader flow, for example).

Liquidity re-equilibration (R5). Every published regression is proximately the frozen one-lag maker’s reversion dominating once the spiral is damped, so the ceiling could in part reflect hand-set balance. Re-tuning the maker strength k_M jointly with each rationality layer (six candidates, tuning seeds 900–903, least-re-equilibration tie-break): the anticipator arm keeps the frozen $k_M = 0.30$ and scores 5/10 on evaluation with the control’s failure set; the full-absorption arm selects $k_M = 0.10$ and scores 2/10, passing only efficiency and the one-clock event, while its raw statistic collapses to a median of 0.0016 bits (0.0007–0.0032 per seed, still

Config	raw bits	whitened bits	sign-alone MI	$AC_1(r)$	whitened sig.
FCVM (control)	0.0184	0.0197	0.0101	+0.036	8/8
$K=1$	0.0200	0.0112	0.0147	-0.075	8/8
$K=5$ (frozen)	0.0257	0.0038	0.0220	-0.247	8/8
FCVM+Q($\lambda_Q=0.5$)	0.0180	0.0033	0.0157	-0.160	7/8
FCVM+Q($\lambda_Q=1.0$)	0.0237	0.0036	0.0206	-0.256	8/8

Table 7. The E9 attribution (medians over the 8 evaluation seeds; worlds regenerated and asserted against the stored per-seed values before any new estimator was read). Raw bits: the published E9 statistic. Whitened bits: the same statistic recomputed on AR(1)-whitened returns — the genuine nonlinear sign leak. Sign-alone MI: the information carried by the $\text{sign}(r_t)$ feature alone — the linear-reversal channel. Every intervention moves information from the whitened column into the sign-alone column while driving $AC_1(r)$ negative: conversion, not creation. The whitened leak stays significant on 8 of 8 seeds in every row except half skew (7 of 8).

E9-significant). That world is the program’s cleanest joint-production exhibit: adapted liquidity plus exact absorption comes within a hair of sign efficiency, and the price is every amplitude fact at once. The ceiling survives re-equilibration of the liquidity parameter; the claim is scoped to the one parameter tested.

Threshold robustness (R6). The battery’s 12 numeric thresholds were perturbed one at a time by $\pm 10\%$ and $\pm 20\%$ of their magnitude (48 perturbed scorings; E9 is a permutation test and has no numeric threshold to perturb), with every score recomputed from per-seed statistics: the rescorer reproduces all fifteen published configuration scores and SPY’s before any perturbed score is read, with the eight configurations lacking archived per-seed statistics reproduced by one asserted deterministic rerun each. Under all 48 perturbations, FCVM remains the maximum simulated score and its failure set never changes; GBM stays 3/10 everywhere; SPY stays 10/10 at $\pm 10\%$ and drops to 9/10 only for the E8 ratio at -20% . The headline rows are frozen at $\pm 10\%$; the knife edges are confined to non-headline configurations (FCVM+Q($\lambda_Q=0.5$), FCVMH, and the frozen $K=5$ stack) and are enumerated, with the full $\pm 20\%$ bands, in the archived results (Appendix C). The integer comparative structure — the 5/10 ceiling and both calibration anchors — is not a threshold artifact.

Composed, R1–R6 replace the natural reading of Experiments II and III with a sharper one. What died: “deeper anticipation grows the sign leak” (a strength axis, R1) and “the market re-creates the information in price space” (induced reversal, R2 — well-powered as a raw effect at 32 seeds, R3). What survives is stronger than what was lost: added rationality *converts* a genuine nonlinear multi-day sign leak into one-lag microstructure reversal — absorbing most of it, never closing it — and pays for the conversion with the amplitude facts, which are flow-generated (R4) and cannot be rescued by re-equilibrating liquidity provision (R5); the ceiling and both calibration anchors are threshold-robust (R6).

6.6 Composition: the structural claim

The three refutations and their attribution compose. In this model class, the forecastable flow component and the asymmetric amplitude facts are the same object:

1. an agent who *trades* against the forecastable flow adds its own state-driven flow: it absorbs




Figure 7. Conversion, not creation. Raw direction bits (black circles; the E9 statistic) rise from the control to the strong-intervention configurations, while the AR(1)-whitened bits (red squares) — the genuinely nonlinear component of the leak — fall monotonically and remain significantly positive in every configuration. Small markers show the eight evaluation seeds; large markers connect the medians. The divergence of the two series is the attribution result of Section 6.5: added rationality converts the leak into first-order reversal rather than removing it.

part of the genuine leak and manufactures one-lag reversal in its place, leaving the raw statistic unchanged (Experiment I with Section 6.5);

2. stacking such agents to the fixed point of mutual anticipation is algebraically — and, at matched strength, measurably — one stronger such agent: depth adds nothing, and the added strength converts more of the leak while the model-space forecastable flow contracts to 0.237 (Experiment II);
3. a price-setter who absorbs the flow’s impact exactly — the rational-expectations limit for this channel, verified to the identity — absorbs most of the genuine leak (whitened bits $0.0197 \rightarrow 0.0036$) but silences none of it, and deletes the de-leveraging dynamics that generate the tails, the leverage effect, and the crash asymmetry, scoring 1/10 (Experiment III).

Sign efficiency and the amplitude facts are not separate dials: in a market moved by state-dependent flows they are jointly produced — a statement measured three independent ways, including the re-equilibrated 2/10 world in which near-closure of the sign channel costs every amplitude fact at once, and scoped by the $t(3)$ probe to flow-generated amplitude structure — and no added layer of rationality — trading, iterated, or price-setting — supplies the five missing events (E3, E4, E7, E8, E9). What the minimal market lacks is not a smarter agent but a different kind of memory: the long-memory clock and the coupling arrow live upstream of the flows, in whatever generates them.




Figure 8. The leverage effect $\text{corr}(r_t, \text{RV}_{t+1..10})$ at three microstructure vertices. Filled points: cohort estimates ± 1 s.d. Open circles: the 13 individual FX crosses and 10 individual coins. Equities (four universes: US, Japan, Europe, Asia-EM) sit at -0.030 to -0.047 ; FX at $+0.005 \pm 0.009$ (statistically zero, $z = 0.56$); crypto at $+0.031 \pm 0.016$. The equity–FX separation is resolved ($z = 3.89$); the FX–crypto separation is not ($z = 1.44$).

7 External validity: the leverage effect across market structures

The model mechanism makes the amplitude channel a property of the *flow structure*: the leverage effect in the minimal market exists because a one-directional, state-dependent de-leveraging flow exists, and disappears when that flow’s impact is removed (Figure 5). A weak but testable empirical counterpart follows: across real asset classes, the leverage effect should follow the microstructure that does or does not host such flows.

Table 8 and Figure 8 give the measurement, using an estimator battery identical across the three vertices (realized variance from daily OHLC ranges via Garman–Klass (Garman and Klass, 1980)). The classical readings of the equity effect — firm capital structure (Christie, 1982) and volatility feedback through the equity risk premium (Bekaert and Wu, 2000) — are equity-specific mechanisms, which is exactly what makes the cross-venue comparison informative: a currency is not a levered claim and carries no equity premium, so whatever asymmetry survives outside equities must come from the venue’s flow structure. Equity markets — where financial leverage and institutional de-risking cascades live — show the textbook negative effect in all four universes (-0.030 to -0.047). FX — an institutional dealer market hosting neither de-risking cascades nor a dominant retail-momentum dynamic — is statistically zero ($+0.005$, z vs. zero = 0.56 ; 7 of 13 pairs positive, the symmetric scatter a true zero predicts). Crypto — no financial leverage, no institutional de-risking, a retail dynamic in which rallies rather than crashes spike volatility — sits at $+0.031$, with 8 of 10 coins individually positive and only BTC (-0.043) and ETH (-0.019) retaining the equity-style sign.

A distinction between what is statistically certified and what is suggested by point estimates is required here: the equity edge is resolved (one-sided $z = 3.89$ against FX, $z = 4.06$ against crypto), and FX is decisively zero; but the FX–crypto edge is $z = 1.44$, inside noise, and crypto

	equity cohort	FX	crypto
leverage effect	-0.0405 ± 0.0079	$+0.005 \pm 0.009$	$+0.031 \pm 0.016$
z vs. zero	-5.1	0.56	1.97
z vs. FX	3.89	—	1.44
instruments positive	0 of 4 markets	7 of 13 pairs	8 of 10 coins

Table 8. The leverage triangle. Equity cohort: four universes (US, Japan, Europe, Asia-EM), 2010–2026. FX: 13 liquid crosses, 2010–2026. Crypto: 10 major coins, 2017–2026. Each z vs. zero is the cohort mean divided by its cross-sectional dispersion (the $\pm$ values in the first row). The ordering $-0.041 < +0.005 < +0.031$ is monotone in point estimates; the equity edge is statistically certified, the FX–crypto edge is inside one-sided noise.

itself clears zero at only $z \approx 2$ — so “crypto genuinely positive above an FX zero” is an ordering of point estimates, not a certified separation.

That edge is not under-analyzed; it is underpowered, and the bound is arithmetic. Certifying $z \geq 2$ at the current point separation requires the crypto sampling SD to fall from 0.0158 below 0.0099, a $1.6\times$ reduction; the SDs are time-block bootstrap SDs scaling as $1/\sqrt{T}$, so history alone would need ≈ 22 years of daily crypto data against the 8.6 that exist (major-coin daily history begins in late 2017). Nor can the cross-section substitute for T : a registered widening of the universe from 10 to 17 majors — every liquid 2017-era coin passing the pre-declared real-range and history gates — left the sampling SD essentially unchanged ($0.0158 \rightarrow 0.0161$) while the pooled estimate moved to $+0.026$ (14 of 17 coins positive) and the z against FX to 1.16; coins share their volatility events, so widening adds instruments but almost no independent information. (The equity–crypto separation, by contrast, survives the widening at $z = 3.71$.) At these point estimates the FX–crypto edge becomes certifiable only by the accrual of history — not by more coins, more crosses, or a different test — and we report it as what it is: an ordering the data cannot yet promote to a separation.

The residual per-pair scatter in FX (Figure 9) lies on the flight-to-quality axis: the three most negative crosses are all yen pairs (AUDJPY -0.046 , GBPJPY -0.024 , EURJPY -0.022 — yen strength is risk-off, so “down $\rightarrow$ vol up” is flight-to-quality mimicking the leverage effect), and the most positive is USDMXN ($+0.075$, the same flow with the opposite quote orientation).

The reading is licensed by a sign-recovery gate: on synthetic worlds with a known leverage sign (equity-negative, inverted-positive, symmetric-zero), the estimator recovers each with wide separation and its spurious-leverage bound is $|0.04|$, so a zero is read as a zero and an inversion as an inversion.

Finally, the triangle’s epistemic status is a stated boundary of the paper’s claims, and it is worth being precise about what would move it. The triangle is *consistent with* the thesis — the amplitude channel follows the flow structure of the venue, negative exactly where forced de-risking lives, zero in the institutional market without it, and (in point estimate) positive where an opposite retail dynamic dominates — but consistency with three ordered points is weak evidence by construction: any monotone story over the same venue ordering fits them. What would count as proof is a *discriminating* cross-venue prediction, and the mechanism makes one: a venue whose flow structure changes should see its leverage sign migrate. Crypto’s institutionalization is the live case — the mechanism predicts its leverage effect drifts toward the equity sign as leveraged institutional flow displaces retail momentum, and BTC (-0.043) and ETH (-0.019), the two most institutionally held coins, already carry the equity-style sign




Figure 9. The leverage effect per instrument, grouped by venue class: the four equity-universe estimates, the 13 FX crosses, and the 10 registered cryptocurrencies, each sorted within its group. Every equity universe is negative; the FX crosses scatter symmetrically about zero along the flight-to-quality axis; the coins are predominantly positive, with BTC and ETH the exceptions.

in the registered universe (joined by XMR in the widened one) — an in-sample observation that is consistent but proves nothing. Delivering that prediction as a certified test needs either the history the power bound above shows does not yet exist, or flow decompositions (who is trading, not just what the price did) that public daily data does not carry. So the claim stays scoped: the triangle is an empirical counterpart, not an empirical test, and no claim is made that the minimal market’s specific mechanism generates the real effect.

8 Discussion and limitations

8.1 Scope of the claims

The structural claim of Section 6.6 is scoped, and the scoping is part of the result. The following are explicitly *not* claimed: (i) anything about real markets’ generating process; (ii) anything beyond this model class — no learning agents, no heterogeneous information, no strategic equilibrium or rational-expectations solution concept; (iii) optimality of the agent designs — each is gate-verified to *function*, not to be optimal; (iv) any creation of sign information by the interventions — the raw E9 rises of Experiments II and III are measured efficiency breakage (induced one-lag reversal), and the contrary reading is withdrawn in Section 6.5; (v) joint production beyond flow-generated amplitude structure — the $t(3)$ probe leaves the absorption-proof exogenous-amplitude case untested.

“Structural” here means: invariant over the searched class under the pre-registered budgets, with each searched member licensed by a mechanism gate proving it expresses its hypothesis. The claim would be falsified by an agent class outside the three — learning agents that condition on more than the public volatility state, or information-heterogeneous agents whose inventories do not ride a common public state — breaking the ceiling under the same battery and the same protocol. That is an open research direction, and we name it as such rather than hedge

the claim through it.

8.2 The closure clause as a stopping rule

Experiment III’s registration contained, alongside its hypotheses and budget, a closure clause: this is the third and final attempt on the expectation line; if the primary hypothesis is refuted at this budget, the 5/10 ceiling is declared structural for flow-based minimal markets with boundedly-rational agents of every class here considered, and the line closes — no fourth experiment. The clause is not rhetorical; it is the instrument that converts three negative results into one positive claim. An open-ended sequence of failed attempts supports only “we have not yet found the right agent.” A pre-committed final attempt, refuted, supports a stronger statement about the class that was searched — precisely because the search was budgeted, the budget was spent, and the stopping rule was declared before the last result was known. The referee program of Section 6.5 was registered and run after the clause fired, and stayed within it: R1–R6 are attribution and robustness measurements of existing results, not a fourth attempt on the ceiling, and the clause remains in force.

8.3 Limitations

Four limitations are measured rather than merely acknowledged. First, the battery’s absolute anchor is a US-equity statement: the registered multi-index audit (Appendix B) shows the joint 10/10 calibration extends across the large-cap US complex but not to non-US single indices, which score 5–7/10 with the clock-structure and sign events recurring among the misses; every conclusion in this paper is comparative — same events, same seeds — and is unaffected, but absolute realism claims for non-US venues are not licensed. Second, the battery scores rest on eight evaluation seeds; the seed-bootstrap standard errors reported in every experiment table bound the resulting score uncertainty at half an event, the two comparisons that carry the mechanism reading were extended to 32 seeds under a registered budget (Section 6.5), and the three refutations additionally rest on pre-registered criteria and paired seeds rather than on score margins. Third, the joint-production claim is scoped to flow-generated amplitude structure: the registered $t(3)$ probe shows the fundamentalist channel filters exogenous tails (kurtosis 3.06 under full absorption even with $t(3)$ fundamentals), so this design cannot test whether an absorption-proof exogenous amplitude source would decouple the two properties. Fourth, the FX–crypto edge of the leverage triangle is underpowered, and Section 7 derives the bound: certification requires a history that will not exist for roughly another decade, and widening the cross-section demonstrably does not substitute. Each limitation is stated with the measurement that establishes it, which is the form of limitation a pre-registered program can offer.

9 Conclusion

A minimal flow-based market reproduces half of the stylized-fact battery — efficiency, fat tails, volatility clustering amplitude, the leverage effect, crash asymmetry — and no rationality mechanism of three pre-registered kinds moves it past 5/10. A registered attribution program establishes what the failed mechanisms do instead: each converts a genuine nonlinear multi-day sign leak into one-lag microstructure reversal, absorbing most of it — whitened information falls

from 0.0197 bits in the control to 0.0036 under exact absorption — and closing none of it, while the amplitude structure is deleted alongside the absorption. The refutations identify the reason: the forecastable component of flow that keeps the sign channel open is the same object that generates the heavy-tailed, asymmetric amplitude structure, so removing one removes the other. In this model class, sign efficiency and the amplitude facts are jointly produced — a claim structural to flow-generated amplitude structure, robust to liquidity re-equilibration (where near-closure of the sign channel costs every amplitude fact at once), and not an artifact of the battery’s thresholds. The missing events — long memory, the coupling arrow, information-free signs — appear to live upstream of the flows, in whatever generates them, and breaking the ceiling within the protocol would require an agent class outside the three tested here. The cross-venue leverage evidence is consistent with the mechanism’s central implication — the amplitude channel follows the venue’s flow structure — while remaining, as measured, an ordering rather than a certified separation.

Code and data availability

Every number in this paper is regenerable from the public repository. The research JSONs consumed by the figures and tables — under `research/`: `decathlon.json`, `decathlon2.json`, `decathlon3.json`, `decathlon4.json`, `robustness.json`, `fx.json`, `crypto.json`, `crypto_wide.json`, `battery_audit.json`, and `score_se.json` — are versioned in the repository and rebuilt by `run_research.py`; the battery code is identical for real data and simulated output; the gate suite runs in continuous integration on every commit, including the byte-identity pins that protect every published row in this paper; and every figure is generated from the research JSONs by `docs/paper/make_figures.py` with no hand-entered values. The pre-registration documents, including the amendments and the closure clause quoted here, are versioned in `docs/design/`.

References

- Tobias Adrian and Hyun Song Shin. Liquidity and leverage. *Journal of Financial Intermediation*, 19(3):418–437, 2010. doi: 10.1016/j.jfi.2008.12.002.
- David H. Bailey and Marcos López de Prado. The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40(5):94–107, 2014. doi: 10.3905/jpm.2014.40.5.094.
- David H. Bailey, Jonathan Borwein, Marcos López de Prado, and Qiji Jim Zhu. The probability of backtest overfitting. *Journal of Computational Finance*, 20(4):39–69, 2017. doi: 10.21314/JCF.2016.322.
- Geert Bekaert and Guojun Wu. Asymmetric volatility and risk in equity markets. *Review of Financial Studies*, 13(1):1–42, 2000. doi: 10.1093/rfs/13.1.1.
- Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986. doi: 10.1016/0304-4076(86)90063-1.

- Jean-Philippe Bouchaud, Andrew Matacz, and Marc Potters. Leverage effect in financial markets: The retarded volatility model. *Physical Review Letters*, 87(22):228701, 2001. doi: 10.1103/PhysRevLett.87.228701.
- Jean-Philippe Bouchaud, J. Doyne Farmer, and Fabrizio Lillo. How markets slowly digest changes in supply and demand. In Thorsten Hens and Klaus Reiner Schenk-Hoppé, editors, *Handbook of Financial Markets: Dynamics and Evolution*. North-Holland, 2009. doi: 10.1016/b978-012374258-2.50006-3.
- Jean-Philippe Bouchaud, Julius Bonart, Jonathan Donier, and Martin Gould. *Trades, Quotes and Prices: Financial Markets Under the Microscope*. Cambridge University Press, 2018. doi: 10.1017/9781316659335.
- William A. Brock and Cars H. Hommes. A rational route to randomness. *Econometrica*, 65(5):1059–1095, 1997. doi: 10.2307/2171879.
- William A. Brock and Cars H. Hommes. Heterogeneous beliefs and routes to chaos in a simple asset pricing model. *Journal of Economic Dynamics and Control*, 22(8–9):1235–1274, 1998. doi: 10.1016/S0165-1889(98)00011-6.
- Markus K. Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *Review of Financial Studies*, 22(6):2201–2238, 2009. doi: 10.1093/rfs/hhn098.
- Damien Challet and Yi-Cheng Zhang. Emergence of cooperation and organization in an evolutionary game. *Physica A: Statistical Mechanics and its Applications*, 246(3–4):407–418, 1997. doi: 10.1016/S0378-4371(97)00419-6.
- Damien Challet, Matteo Marsili, and Yi-Cheng Zhang. *Minority Games: Interacting Agents in Financial Markets*. Oxford University Press, 2004. doi: 10.1093/oso/9780198566403.001.0001.
- Carl Chiarella, Roberto Dieci, and Xue-Zhong He. Heterogeneity, market mechanisms, and asset price dynamics. In Thorsten Hens and Klaus Reiner Schenk-Hoppé, editors, *Handbook of Financial Markets: Dynamics and Evolution*. North-Holland, 2009. doi: 10.1016/b978-012374258-2.50009-9.
- Andrew A. Christie. The stochastic behavior of common stock variances: Value, leverage and interest rate effects. *Journal of Financial Economics*, 10(4):407–432, 1982. doi: 10.1016/0304-405X(82)90018-6.
- Rama Cont. Empirical properties of asset returns: Stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001. doi: 10.1080/713665670.
- Rama Cont and Jean-Philippe Bouchaud. Herd behavior and aggregate fluctuations in financial markets. *Macroeconomic Dynamics*, 4(2):170–196, 2000. doi: 10.1017/S1365100500015029.
- Zhuanxin Ding, Clive W. J. Granger, and Robert F. Engle. A long memory property of stock market returns and a new model. *Journal of Empirical Finance*, 1(1):83–106, 1993. doi: 10.1016/0927-5398(93)90006-D.
- Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4):987–1007, 1982. doi: 10.2307/1912773.

- Eugene F. Fama. Efficient capital markets: A review of theory and empirical work. *Journal of Finance*, 25(2):383–417, 1970. doi: 10.2307/2325486.
- J. Doyne Farmer and Shareen Joshi. The price dynamics of common trading strategies. *Journal of Economic Behavior & Organization*, 49(2):149–171, 2002. doi: 10.1016/S0167-2681(02)00065-3.
- Vladimir Filimonov and Didier Sornette. Quantifying reflexivity in financial markets: Toward a prediction of flash crashes. *Physical Review E*, 85(5):056108, 2012. doi: 10.1103/PhysRevE.85.056108.
- Mark B. Garman and Michael J. Klass. On the estimation of security price volatilities from historical data. *Journal of Business*, 53(1):67–78, 1980. doi: 10.1086/296072.
- Jim Gatheral, Thibault Jaisson, and Mathieu Rosenbaum. Volatility is rough. *Quantitative Finance*, 18(6):933–949, 2018. doi: 10.1080/14697688.2017.1393551.
- Irene Giardina and Jean-Philippe Bouchaud. Bubbles, crashes and intermittency in agent based market models. *European Physical Journal B*, 31(3):421–437, 2003. doi: 10.1140/epjb/e2003-00050-6.
- Lawrence R. Glosten, Ravi Jagannathan, and David E. Runkle. On the relation between the expected value and the volatility of the nominal excess return on stocks. *Journal of Finance*, 48(5):1779–1801, 1993. doi: 10.1111/j.1540-6261.1993.tb05128.x.
- Stephen J. Hardiman, Nicolas Bercot, and Jean-Philippe Bouchaud. Critical reflexivity in financial markets: A Hawkes process analysis. *European Physical Journal B*, 86(10):442, 2013. doi: 10.1140/epjb/e2013-40107-3.
- Campbell R. Harvey, Yan Liu, and Heqing Zhu. ... and the cross-section of expected returns. *Review of Financial Studies*, 29(1):5–68, 2016. doi: 10.1093/rfs/hhv059.
- Cars H. Hommes. Heterogeneous agent models in economics and finance. In Leigh Tesfatsion and Kenneth L. Judd, editors, *Handbook of Computational Economics*, volume 2, pages 1109–1186. North-Holland, 2006. doi: 10.1016/S1574-0021(05)02023-X.
- Alexander Kraskov, Harald Stögbauer, and Peter Grassberger. Estimating mutual information. *Physical Review E*, 69(6):066138, 2004. doi: 10.1103/PhysRevE.69.066138.
- Blake LeBaron. Agent-based computational finance: Suggested readings and early research. *Journal of Economic Dynamics and Control*, 24(5–7):679–702, 2000. doi: 10.1016/S0165-1889(99)00022-6.
- Blake LeBaron. Agent-based computational finance. In Leigh Tesfatsion and Kenneth L. Judd, editors, *Handbook of Computational Economics*, volume 2, pages 1187–1233. North-Holland, 2006. doi: 10.1016/S1574-0021(05)02024-1.
- Thomas Lux and Michele Marchesi. Scaling and criticality in a stochastic multi-agent model of a financial market. *Nature*, 397(6719):498–500, 1999. doi: 10.1038/17290.
- Benoit Mandelbrot. The variation of certain speculative prices. *Journal of Business*, 36(4):394–419, 1963. doi: 10.1086/294632.

- Liam Paninski. Estimation of entropy and mutual information. *Neural Computation*, 15(6): 1191–1253, 2003. doi: 10.1162/089976603321780272.
- Paul A. Samuelson. Proof that properly anticipated prices fluctuate randomly. *Industrial Management Review*, 6(2):41–49, 1965. doi: 10.1142/9789814566926_0002. DOI refers to the 2015 reprint in *The World Scientific Handbook of Futures Markets*.
- Andrew Urquhart. The inefficiency of Bitcoin. *Economics Letters*, 148:80–82, 2016. doi: 10.1016/j.econlet.2016.09.019.
- Wei Zhang, Pengfei Wang, Xiao Li, and Dehua Shen. Some stylized facts of the cryptocurrency market. *Applied Economics*, 50(55):5950–5965, 2018. doi: 10.1080/00036846.2018.1488076.
- Gilles Zumbach. Time reversal invariance in finance. *Quantitative Finance*, 9(5):505–515, 2009. doi: 10.1080/14697680802616712.

A The battery: exact criteria and calibration statistics

Table 9 states each event’s statistic and pass region exactly as implemented in the repository (`kronos/decathlon.py`, function `battery`; the thresholds were fixed in the DESIGN8 registration and its calibration-phase amendments), and Table 10 gives the full statistic vector at both calibration ends and at the flow-only ceiling. Notation: r_t are daily returns; RV_t is the 5-day rolling mean of r^2 (the close-only volatility clock), $\sigma_t = \sqrt{RV_t}$ (written σ_{5d} in Table 2), and $z_t = r_t/\sigma_t$ the deformed (one-clock) return. The weekly clock is built from *non-overlapping* 5-day means w_j of $\log r^2$ — raw weekly differences carry an MA(1) artifact on which even GBM’s $|\Delta w|$ autocorrelates, so E4 uses the level and E7 uses AR(1) innovations $u_j = w_{j+1} - \bar{w}(1 - \phi) - \phi w_j$ with ϕ the fitted AR(1) coefficient. $AC_k(\cdot)$ is the lag- k autocorrelation. A configuration passes an event if a majority of its 8 evaluation seeds pass.

Two calibration facts anchor the thresholds. GBM’s weekly-clock innovation skew (-0.651) is the $\log\text{-}\chi^2$ noise floor E7 is calibrated against; real volatility’s up-jumps cancel most of it (SPY -0.221). And the E4 level-autocorrelation criterion is what a GJR-GARCH process fails (clock AC_8 median 0.03, maximum 0.09, over the battery’s eight evaluation seeds of the gate suite’s GJR world — every seed below the 0.12 bar): exponential memory decays too fast at eight weeks, which is why E4 is the battery’s long-memory event rather than any Hurst point estimate (the close-only Hurst proxy is biased and is recorded as informational only).

B The multi-index calibration audit

The joint calibration anchor of Section 4.1 is one index. A registered audit (fixed list, declared before any score was read) ran the identical close-only battery on six further real indices: QQQ, DIA, IWM (2010–2026), and the three locally-listed index ETFs used by the research program’s transfer study — 1306.T (TOPIX), EXW1.DE (EURO STOXX 50), 2800.HK (Hang Seng tracker). The registered prediction — every index $\geq 9/10$ — *failed*, and the failure is itself a measurement (Table 11): the anchor extends across the large-cap US complex (DIA joins SPY at 10/10; QQQ and IWM drop single events), but non-US single indices score 5–7/10, missing the clock-structure and sign events (E7, E8, E9 recur) — consistent with the transfer study’s

#	Statistic	Pass region
E1	$AC_1(r)$	$[-0.15, 0.05]$
E2	$kurtosis(r)$	$[4.5, 40]$
E3	$AC_1(r)$ and $\text{mean}_{k=5..20} AC_k(r)$	≥ 0.12 and ≥ 0.05
E4	$AC_8(w)$ (weekly clock level)	≥ 0.12
E5	$\text{mean}_{\tau=1..10} \text{corr}(r_t, RV_{t+\tau})$	≤ -0.03
E6	$kurtosis(z)$	≤ 5
E7	$\text{skew}(u)$ (weekly clock AR(1) innovations)	≥ -0.35
E8	$EP(r), EP(z)$ (entropy production, bits)	$EP(r)$ significant and $EP(z) \leq 0.75 EP(r)$
E9	direction bits $I(\text{features}; \text{sign } r_{t+1})$	$\leq 95\text{th percentile of 120 shuffles}$
E10	$\#\{r < -2.5s\} / \#\{r > +2.5s\}, s = \text{sd}(r)$	≥ 1.25

Table 9. Exact pass criteria, as implemented. E7’s floor is calibrated against the $\log\text{-}\chi^2$ noise skew of the weekly clock (GBM innovations skew ≈ -0.65 ; real SPY -0.22). E8’s entropy production symbolizes the series into 3 bins, measures 3-gram irreversibility, and tests against 120 block-reversal surrogates (random block offsets, blocks reversed with probability 1/2). E9’s features are the Cartesian product of three tape observables — $\text{sign}(r_t)$, the sign of 21-day momentum, and the volatility tercile — scored by discrete plug-in mutual information with Miller–Madow bias correction (see Paninski, 2003) against a 120-permutation null; the program’s continuous information channel uses the KSG estimator (Kraskov et al., 2004), but E9’s features are discrete by construction. The volatility tercile is cut on full-sample quantiles: E9 is a measurement of information content, computed identically on real and simulated series, not a trading rule.

finding that mechanism transfers across venues while calibration values are local. A hygiene probe recorded in the registration rules out the union-calendar padding of the foreign panels as the cause (native-calendar re-runs score the same or lower). The battery’s absolute anchor is therefore a US-equity statement; all conclusions in this paper are comparative and unaffected. Figure 10 shows the full event-by-index pass/fail matrix.

C Robustness material

Score standard errors. The seed-bootstrap standard-error method of Section 4.1 was applied to every configuration published in this paper (fifteen in total, spanning the ablation ladder and all three experiments). The bootstrap reproduces every published majority-vote score exactly before resampling, and the resulting SEs range from 0.00 to 0.49 events (Figure 11); configurations whose per-seed event counts sit at a majority knife-edge (F, the frozen $K=5$ stack, FCVMH) carry the largest SEs, while the headline comparisons (FCVM, FCVM+A, the tuned configurations) are tight.

Threshold perturbation bands. The full $\pm 20\%$ single-threshold score bands for all fifteen published configurations, SPY’s band, and the named knife-edge (configuration, threshold) pairs of Section 6.5 are archived in `research/robustness.json (threshold.robustness)`; the per-perturbation records include every configuration’s score and FCVM’s failure set under each of the 48 perturbations. At $\pm 20\%$ the widest bands belong to the knife-edge configurations (the frozen $K=5$ stack spans 2–4 events); every headline row stays at exactly 5.

Statistic	Event	SPY (real)	GBM (G)	FCVM
$AC_1(r)$	E1	-0.103	0.001	0.036
kurtosis(r)	E2	15.09	3.04	8.79
$AC_1(r)$	E3	0.307	-0.001	0.377
mean $AC_{5..20}(r)$	E3	0.218	0.002	-0.050
$AC_8(w)$	E4	0.186	-0.018	-0.003
leverage	E5	-0.113	0.007	-0.125
kurtosis(z)	E6	2.33	2.13	2.20
skew(u)	E7	-0.221	-0.651	-0.408
EP(r) bits	E8	0.0211	0.0002	0.0929
EP(z) bits	E8	0.0141	0.0006	0.1392
direction bits	E9	0.0029	0.0001	0.0184
tail asymmetry	E10	1.56	0.91	39.96

Table 10. The full statistic vector at the battery’s calibration ends (SPY 10/10, GBM 3/10) and at the flow-only ceiling (FCVM, 5/10); GBM and FCVM entries are medians over the 8 evaluation seeds. The table shows *how* each score decomposes: GBM passes only E1, E6, E9; FCVM matches SPY’s magnitudes on E1/E2/E5/E10 but has no clock memory ($AC_8(w) \approx 0$ against SPY’s 0.186), fails E8 on the ratio criterion ($EP(z) > 0.75 EP(r)$: deformation does not reduce its entropy production), and leaks $6\times$ SPY’s direction bits. Its E7 skew (-0.408) sits marginally below the -0.35 floor. SPY’s $EP(z)$ is significant but 33% below $EP(r)$, meeting the ratio criterion.

The widened cryptocurrency universe. The registered widening of the crypto panel from 10 to 17 majors (Section 3.3) supports the power analysis of Section 7: it leaves the cross-sectional sampling SD essentially unchanged while moving the pooled leverage estimate to $+0.026$ (14 of 17 coins positive), which is the measured demonstration that the FX–crypto separation is history-bound rather than cross-section-bound.

	E1	E2	E3	E4	E5	E6	E7	E8	E9	E10	
SPY (anchor)	■	■	■	■	■	■	■	■	■	■	10/10
QQQ	■	■	■	■	■	■	□	□	■	■	8/10
DIA	■	■	■	■	■	■	■	■	■	■	10/10
IWM	■	■	■	■	■	■	■	□	■	□	8/10
1306.T	■	■	■	□	■	■	□	□	□	■	6/10
EXW1.DE	■	■	■	■	■	■	□	□	□	■	7/10
2800.HK	■	■	□	□	■	■	□	■	□	□	5/10

Figure 10. The multi-index calibration audit as an event-by-index matrix (filled = pass). The clock-structure and sign events (E7, E8, E9) recur among the non-US misses.

Index	Venue	Score	Failed events
SPY (anchor)	US	10/10	—
DIA	US	10/10	—
QQQ	US	8/10	E7 E8
IWM	US	8/10	E8 E10
1306.T	Japan	6/10	E4 E7 E8 E9
EXW1.DE	Europe	7/10	E7 E8 E9
2800.HK	Hong Kong	5/10	E3 E4 E7 E9 E10

Table 11. The multi-index calibration audit (registered list, no post-hoc swaps; misses reported by name). US large-cap indices reproduce the anchor; non-US single indices do not — the battery’s thresholds are US-calibrated, so its absolute scores certify realism only against US equity data, which is how the paper uses them.




Figure 11. Battery scores with seed-bootstrap standard errors for every configuration published in this paper. The dotted guide marks the 5/10 flow-only ceiling. No SE exceeds half an event.